

3. Why LLMs Behave Differently in Production Example

Intro

This example explains why the same prompt can behave differently in testing and production. The **LLM INQUISITOR METHODOLOGY** describes how background conditions affect behaviour and how evaluators decide whether to continue or stop based on whether further probing will add useful information. This example gives a simple demonstration that does not require reading the full method first.

Step 1: Clean start

You begin with a simple task. The model is fresh, the conversation is short, and it performs well.

Step 2: Conversation gets longer

As you continue:

- the model has more context to juggle
- earlier details get buried
- it starts simplifying or drifting

Even if your instructions stay simple, the internal pressure increases.

Step 3: Same prompt, different behaviour

You give the same prompt again later in the conversation. This time the model:

- forgets details
- contradicts earlier answers
- oversimplifies
- mixes in unrelated information

Why this happens

In testing, prompts are clean and isolated. In production, prompts arrive:

- after long conversations
- with more context
- with more constraints
- with accumulated mistakes

So the same prompt behaves differently.

What this shows

Inconsistency is not random. It is tied to the conditions the model is operating under. The evaluator uses this to decide whether further probing will add useful information or whether the behaviour has already been sufficiently characterised.